

In [1]:

```
# =====
# CELL 1: IMPORTS AND SETUP
# =====
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.cluster import KMeans
from sklearn.mixture import GaussianMixture
from sklearn.decomposition import PCA, FastICA
from sklearn.random_projection import GaussianRandomProjection
from sklearn.metrics import silhouette_score
from sklearn.neural_network import MLPClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, f1_score
from scipy.stats import kurtosis
import time
import warnings
warnings.filterwarnings('ignore')

np.random.seed(42)
plt.style.use('seaborn-v0_8')

print("=*70")
print("UNSUPERVISED LEARNING AND DIMENSIONALITY REDUCTION")
print("CS7641 Machine Learning - Summer 2025")
print("Max Matkovski")
print("=*70")
=====
```

UNSUPERVISED LEARNING AND DIMENSIONALITY REDUCTION  
CS7641 Machine Learning - Summer 2025  
Max Matkovski

---

In [2]:

```
# =====
# CELL 2: DATA LOADING AND PREPROCESSING
# =====
print("\n" + "=*70")
print("DATA LOADING AND PREPROCESSING")
print("=*70")

# Load datasets
cancer_df = pd.read_csv('cancer.csv')
bankrupt_df = pd.read_csv('bankrupt.csv')

print("\n--- Cancer Dataset ---")
print(f"Original shape: {cancer_df.shape}")
print(f"Columns: {list(cancer_df.columns)}")

print("\n--- Bankruptcy Dataset ---")
print(f"Original shape: {bankrupt_df.shape}")
print(f"Number of columns: {len(bankrupt_df.columns)}")
```

```

# Clean and preprocess
cancer_df_clean = cancer_df.dropna()
bankrupt_df_clean = bankrupt_df.dropna()

# Select numeric features only
cancer_X = cancer_df_clean.select_dtypes(include=[np.number])
bankrupt_X = bankrupt_df_clean.select_dtypes(include=[np.number])

# Store labels for later use
cancer_labels = cancer_df_clean['Cancer_Type'].values
bankrupt_labels = bankrupt_df_clean['Bankrupt?'].values

print(f"\nCancer - Samples: {cancer_X.shape[0]}, Features: {cancer_X.shape[1]}")
print(f"Cancer features: {list(cancer_X.columns)}")
print(f"\nBankruptcy - Samples: {bankrupt_X.shape[0]}, Features: {bankrupt_X.shape[1]}")

# Standardize features
scaler_cancer = StandardScaler()
scaler_bankrupt = StandardScaler()

cancer_X_scaled = scaler_cancer.fit_transform(cancer_X)
bankrupt_X_scaled = scaler_bankrupt.fit_transform(bankrupt_X)

print("\n✓ Data preprocessing complete")
print(f" Cancer: {cancer_X_scaled.shape}")
print(f" Bankruptcy: {bankrupt_X_scaled.shape}")

```

=====  
DATA LOADING AND PREPROCESSING  
=====

--- Cancer Dataset ---

Original shape: (50000, 15)

Columns: ['Patient\_ID', 'Age', 'Gender', 'Country\_Region', 'Year', 'Genetic\_Risk', 'Air\_Pollution', 'Alcohol\_Use', 'Smoking', 'Obesity\_Level', 'Cancer\_Type', 'Cancer\_Stage', 'Treatment\_Cost\_USD', 'Survival\_Years', 'Target\_Severity\_Score']

--- Bankruptcy Dataset ---

Original shape: (6819, 96)

Number of columns: 96

Cancer - Samples: 50000, Features: 10

Cancer features: ['Age', 'Year', 'Genetic\_Risk', 'Air\_Pollution', 'Alcohol\_Use', 'Smoking', 'Obesity\_Level', 'Treatment\_Cost\_USD', 'Survival\_Years', 'Target\_Severity\_Score']

Bankruptcy - Samples: 6819, Features: 96

✓ Data preprocessing complete

Cancer: (50000, 10)

Bankruptcy: (6819, 96)

In [3]:

```
# =====
# CELL 3: STEP 1 - OPTIMAL K SELECTION
# =====
```

```

print("\n" + "="*70)
print("STEP 1: CLUSTERING ON ORIGINAL DATA")
print("="*70)

def find_optimal_k(X, max_k=10, dataset_name="Dataset"):
    """Find optimal k using elbow method, silhouette score, and BIC"""
    inertias = []
    silhouettes = []
    bics = []
    k_range = range(2, max_k + 1)

    for k in k_range:
        # K-Means
        km = KMeans(n_clusters=k, random_state=42, n_init=10)
        km.fit(X)
        inertias.append(km.inertia_)
        silhouettes.append(silhouette_score(X, km.labels_))

        # GMM
        gm = GaussianMixture(n_components=k, random_state=42, n_init=3)
        gm.fit(X)
        bics.append(gm.bic(X))

    # Create plots
    fig, axes = plt.subplots(1, 3, figsize=(16, 5))

    # Elbow plot
    axes[0].plot(list(k_range), inertias, 'bo-', linewidth=2, markersize=8)
    axes[0].set_xlabel('Number of Clusters (k)', fontsize=12)
    axes[0].set_ylabel('Inertia', fontsize=12)
    axes[0].set_title(f'{dataset_name}: K-Means Elbow Method', fontsize=14,
    axes[0].grid(True, alpha=0.3)

    # Silhouette plot
    axes[1].plot(list(k_range), silhouettes, 'go-', linewidth=2, markersize=8)
    axes[1].set_xlabel('Number of Clusters (k)', fontsize=12)
    axes[1].set_ylabel('Silhouette Score', fontsize=12)
    axes[1].set_title(f'{dataset_name}: Silhouette Analysis', fontsize=14, f
    best_k_sil = list(k_range)[np.argmax(silhouettes)]
    axes[1].axvline(x=best_k_sil, color='r', linestyle='--', linewidth=2,
                    label=f'Best k={best_k_sil}')
    axes[1].legend(fontsize=11)
    axes[1].grid(True, alpha=0.3)

    # BIC plot
    axes[2].plot(list(k_range), bics, 'ro-', linewidth=2, markersize=8)
    axes[2].set_xlabel('Number of Components', fontsize=12)
    axes[2].set_ylabel('BIC Score', fontsize=12)
    axes[2].set_title(f'{dataset_name}: GMM BIC Score', fontsize=14, fontwei
    best_k_bic = list(k_range)[np.argmin(bics)]
    axes[2].axvline(x=best_k_bic, color='b', linestyle='--', linewidth=2,
                    label=f'Best k={best_k_bic}')
    axes[2].legend(fontsize=11)
    axes[2].grid(True, alpha=0.3)

plt.tight_layout()

```

```

plt.savefig(f'{dataset_name.lower()}_optimal_k.png', dpi=150, bbox_inches='tight')
plt.show()

print(f"\n{dataset_name} Optimal k Analysis:")
print(f" Best k by Silhouette: {best_k_sil} (score: {max(silhouettes)}: ")
print(f" Best k by BIC: {best_k_bic}")

return best_k_sil, best_k_bic, silhouettes, bics


# Find optimal k for both datasets
print("\n--- Cancer Dataset: Optimal k Selection ---")
best_k_cancer, _, cancer_sils, cancer_bics = find_optimal_k(cancer_X_scaled,)

print("\n--- Bankruptcy Dataset: Optimal k Selection ---")
best_k_bankrupt, _, bankrupt_sils, bankrupt_bics = find_optimal_k(bankrupt_X_scaled,)

```

=====  
**STEP 1: CLUSTERING ON ORIGINAL DATA**  
=====


--- Cancer Dataset: Optimal k Selection ---



Cancer Optimal k Analysis:

Best k by Silhouette: 2 (score: 0.1198)  
 Best k by BIC: 10

--- Bankruptcy Dataset: Optimal k Selection ---



Bankruptcy Optimal k Analysis:

Best k by Silhouette: 4 (score: 0.1427)  
 Best k by BIC: 5

In [4]:

```

# =====
# CELL 4: PERFORM CLUSTERING ON ORIGINAL DATA
# =====
print("\n--- Clustering on Original Data ---")

```

```

# Use k=3 for consistency
K = 3

def cluster_original(X, k, dataset_name):
    """Perform K-Means and GMM clustering on original data"""
    # K-Means
    km = KMeans(n_clusters=k, random_state=42, n_init=10)
    km_labels = km.fit_predict(X)
    km_sil = silhouette_score(X, km_labels)

    # GMM
    gm = GaussianMixture(n_components=k, random_state=42, n_init=3)
    gm_labels = gm.fit_predict(X)
    gm_sil = silhouette_score(X, gm_labels)

    print(f"\n{dataset_name} Original Data Clustering (k={k}):")
    print(f"  K-Means Silhouette: {km_sil:.4f}")
    print(f"  K-Means Inertia: {km.inertia_:.2f}")
    print(f"  GMM Silhouette: {gm_sil:.4f}")
    print(f"  GMM BIC: {gm.bic(X):.2f}")

    return km_labels, gm_labels, km_sil, gm_sil

# Cluster both datasets
cancer_km_orig, cancer_gm_orig, cancer_km_sil_orig, cancer_gm_sil_orig = \
    cluster_original(cancer_X_scaled, K, "Cancer")

bankrupt_km_orig, bankrupt_gm_orig, bankrupt_km_sil_orig, bankrupt_gm_sil_orig = \
    cluster_original(bankrupt_X_scaled, K, "Bankruptcy")

# Initialize results storage for comprehensive table
all_results = []

# Add original results
all_results.extend([
    {'Dataset': 'Cancer', 'DR_Method': 'None', 'Clustering': 'K-Means',
     'Silhouette': cancer_km_sil_orig},
    {'Dataset': 'Cancer', 'DR_Method': 'None', 'Clustering': 'GMM',
     'Silhouette': cancer_gm_sil_orig},
    {'Dataset': 'Bankruptcy', 'DR_Method': 'None', 'Clustering': 'K-Means',
     'Silhouette': bankrupt_km_sil_orig},
    {'Dataset': 'Bankruptcy', 'DR_Method': 'None', 'Clustering': 'GMM',
     'Silhouette': bankrupt_gm_sil_orig}
])

```

--- Clustering on Original Data ---

Cancer Original Data Clustering (k=3):

K-Means Silhouette: 0.0868  
K-Means Inertia: 411379.10  
GMM Silhouette: 0.0722  
GMM BIC: 806766.82

Bankruptcy Original Data Clustering (k=3):

K-Means Silhouette: 0.1074  
K-Means Inertia: 583013.28  
GMM Silhouette: 0.0338  
GMM BIC: -1611911.63

In [5]:

```
# =====#
# CELL 5: STEP 2a - PCA ANALYSIS WITH SCREE PLOT
# =====#
print("\n" + "="*70)
print("STEP 2: DIMENSIONALITY REDUCTION ANALYSIS")
print("="*70)

def analyze_pca(X, dataset_name):
    """Complete PCA analysis with scree plot"""
    n_comp = min(X.shape[0], X.shape[1])
    pca = PCA(n_components=n_comp, random_state=42)
    pca.fit(X)

    fig, axes = plt.subplots(1, 2, figsize=(14, 5))

    # Scree plot
    eigenvalues = pca.explained_variance_
    n_show = min(20, len(eigenvalues))

    axes[0].plot(range(1, n_show+1), eigenvalues[:n_show], 'bo-', linewidth=2)
    axes[0].set_xlabel('Principal Component', fontsize=12)
    axes[0].set_ylabel('Eigenvalue', fontsize=12)
    axes[0].set_title(f'{dataset_name}: PCA Scree Plot', fontsize=14, fontweight='bold')
    axes[0].grid(True, alpha=0.3)

    # Cumulative variance
    cumvar = np.cumsum(pca.explained_variance_ratio_)
    axes[1].plot(range(1, n_show+1), cumvar[:n_show], 'ro-', linewidth=2, marker='o')
    axes[1].axhline(y=0.95, color='g', linestyle='--', linewidth=2, label='95%')
    axes[1].axhline(y=0.90, color='orange', linestyle='--', linewidth=2, label='90%')
    axes[1].set_xlabel('Number of Components', fontsize=12)
    axes[1].set_ylabel('Cumulative Variance Explained', fontsize=12)
    axes[1].set_title(f'{dataset_name}: Cumulative Variance', fontsize=14, fontweight='bold')
    axes[1].legend(fontsize=11, loc='lower right')
    axes[1].grid(True, alpha=0.3)

    plt.tight_layout()
    plt.savefig(f'{dataset_name.lower()}_pca_scree.png', dpi=150, bbox_inches='tight')
    plt.show()

    n_90 = np.argmax(cumvar >= 0.90) + 1
    n_95 = np.argmax(cumvar >= 0.95) + 1
```

```


print(f"\n{dataset_name} PCA Analysis:")
print(f" Total components: {len(eigenvalues)}")
print(f" Components for 90% variance: {n_90}")
print(f" Components for 95% variance: {n_95}")
print(f" Variance by first 2 PCs: {cumvar[1]*100:.2f}%")
print(f" PC1 variance: {pca.explained_variance_ratio_[0]*100:.2f}%")
print(f" PC2 variance: {pca.explained_variance_ratio_[1]*100:.2f}%")

return pca, cumvar

print("\n--- PCA Analysis ---")
pca_cancer_full, cancer_cumvar = analyze_pca(cancer_X_scaled, "Cancer")
pca_bankrupt_full, bankrupt_cumvar = analyze_pca(bankrupt_X_scaled, "Bankrup
=====
```


## STEP 2: DIMENSIONALITY REDUCTION ANALYSIS

### --- PCA Analysis ---



### Cancer PCA Analysis:

Total components: 10  
Components for 90% variance: 8  
Components for 95% variance: 9  
Variance by first 2 PCs: 30.26%  
PC1 variance: 20.11%  
PC2 variance: 10.15%



```
Bankruptcy PCA Analysis:  
Total components: 96  
Components for 90% variance: 46  
Components for 95% variance: 54  
Variance by first 2 PCs: 20.45%  
PC1 variance: 13.30%  
PC2 variance: 7.14%
```


```
In [6]: # ======  
# CELL 6: STEP 2b - ICA KURTOSIS ANALYSIS  
# ======  
  
def analyze_ica(X, dataset_name, n_components=10):  
    """ICA analysis with kurtosis visualization"""  
    n_comp = min(n_components, X.shape[1])  
  
    ica = FastICA(n_components=n_comp, random_state=42, max_iter=1000)  
    X_ica = ica.fit_transform(X)  
  
    # Calculate kurtosis  
    kurt_vals = [kurtosis(X_ica[:, i], fisher=True) for i in range(n_comp)]  
  
    # Plot  
    fig, ax = plt.subplots(figsize=(12, 5))  
  
    colors = ['darkred' if abs(k) > 1 else 'orange' if abs(k) > 0.5 else 'green'  
              for k in kurt_vals]  
  
    bars = ax.bar(range(1, n_comp+1), kurt_vals, color=colors, alpha=0.7,  
                  edgecolor='black', linewidth=1.5)  
    ax.axhline(y=0, color='blue', linestyle='--', linewidth=2,  
               label='Normal distribution (kurtosis=0)')  
    ax.set_xlabel('Independent Component', fontsize=12)  
    ax.set_ylabel('Kurtosis (Fisher)', fontsize=12)  
    ax.set_title(f'{dataset_name}: ICA Component Kurtosis', fontsize=14, fontweight='bold')  
    ax.legend(fontsize=11)  
    ax.grid(True, alpha=0.3, axis='y')  
  
    # Add value labels  
    for bar, val in zip(bars, kurt_vals):  
        height = bar.get_height()  
        ax.text(bar.get_x() + bar.get_width()/2., height,  
                f'{val:.2f}', ha='center', va='bottom' if height > 0 else 'top')  
  
    plt.tight_layout()  
    plt.savefig(f'{dataset_name.lower()}_ica_kurtosis.png', dpi=150, bbox_inches='tight')  
    plt.show()  
  
    print(f"\n{dataset_name} ICA Kurtosis Analysis:")  
    print(f" Components analyzed: {n_comp}")  
    print(f" Mean |kurtosis|: {np.mean(np.abs(kurt_vals)):.4f}")  
    print(f" Max kurtosis: {np.max(kurt_vals):.4f}")  
    print(f" Components with |kurtosis| > 1: {sum(np.abs(kurt_vals) > 1)}")  
    print(f" Components with |kurtosis| > 0.5: {sum(np.abs(kurt_vals) > 0.5)}")  
  
    return kurt_vals
```

```

print("\n--- ICA Kurtosis Analysis ---")
cancer_kurt = analyze_ica(cancer_X_scaled, "Cancer", n_components=10)
bankrupt_kurt = analyze_ica(bankrupt_X_scaled, "Bankruptcy", n_components=10)


```

--- ICA Kurtosis Analysis ---



Cancer ICA Kurtosis Analysis:

Components analyzed: 10  
 Mean |kurtosis|: 1.2012  
 Max kurtosis: -1.1930  
 Components with |kurtosis| > 1: 10  
 Components with |kurtosis| > 0.5: 10



Bankruptcy ICA Kurtosis Analysis:

Components analyzed: 10  
 Mean |kurtosis|: 981.4846  
 Max kurtosis: 3894.1875  
 Components with |kurtosis| > 1: 10  
 Components with |kurtosis| > 0.5: 10

In [7]:

```

# =====#
# CELL 7: STEP 2c - RANDOM PROJECTION RECONSTRUCTION ERROR
# =====#
def analyze_rp(X, dataset_name, n_runs=10):
    """Random projection reconstruction error analysis"""
    max_comp = min(X.shape[1], 15)
    n_comp_range = range(2, max_comp + 1)

```

```

mean_errors = []
std_errors = []

print(f"Analyzing Random Projection for {dataset_name}...")

for n_comp in n_comp_range:
    errors = []
    for _ in range(n_runs):
        rp = GaussianRandomProjection(n_components=n_comp)
        X_rp = rp.fit_transform(X)
        X_reconstructed = X_rp @ rp.components_
        error = np.mean((X - X_reconstructed) ** 2) / np.mean(X ** 2)
        errors.append(error)

    mean_errors.append(np.mean(errors))
    std_errors.append(np.std(errors))

# Plot
fig, ax = plt.subplots(figsize=(12, 5))
ax.errorbar(list(n_comp_range), mean_errors, yerr=std_errors,
            fmt='go-', capsized=4, capthick=2, linewidth=2, markersize=8,
            label=f'Mean ± Std ({n_runs} runs)', elinewidth=2)
ax.set_xlabel('Number of Random Projections', fontsize=12)
ax.set_ylabel('Normalized Reconstruction Error', fontsize=12)
ax.set_title(f'{dataset_name}: Random Projection Reconstruction Error',
            fontsize=14, fontweight='bold')
ax.legend(fontsize=11)
ax.grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig(f'{dataset_name.lower()}_rp_reconstruction.png', dpi=150, bbox_inches='tight')
plt.show()


print(f"\n{dataset_name} Random Projection Analysis:")
print(f"  Error with 2 components: {mean_errors[0]:.4f} ± {std_errors[0]:.4f}")
print(f"  Error with 5 components: {mean_errors[3]:.4f} ± {std_errors[3]:.4f}")
if len(mean_errors) > 8:
    print(f"  Error with 10 components: {mean_errors[8]:.4f} ± {std_errors[8]:.4f}")

return mean_errors, std_errors

print("\n--- Random Projection Analysis ---")
rp_cancer_err, rp_cancer_std = analyze_rp(cancer_X_scaled, "Cancer", n_runs=n_runs)
rp_bankrupt_err, rp_bankrupt_std = analyze_rp(bankrupt_X_scaled, "Bankruptcy", n_runs=n_runs)

```

--- Random Projection Analysis ---  
Analyzing Random Projection for Cancer...




#### Cancer Random Projection Analysis:

Error with 2 components:  $4.0147 \pm 1.7039$

Error with 5 components:  $2.6694 \pm 1.0391$

Error with 10 components:  $1.0816 \pm 0.2132$

#### Analyzing Random Projection for Bankruptcy...



#### Bankruptcy Random Projection Analysis:

Error with 2 components:  $57.3973 \pm 11.5976$

Error with 5 components:  $18.3481 \pm 2.0723$

Error with 10 components:  $9.9535 \pm 1.5295$

```
In [8]: # =====#
# CELL 8: STEP 3 - CLUSTERING AFTER DR (ALL 12 COMBINATIONS)
# =====#
print("\n" + "="*70)
print("STEP 3: CLUSTERING AFTER DIMENSIONALITY REDUCTION")
print("="*70)

def dr_and_cluster(X, k, dataset_name, n_components=2):
    """Apply DR methods then cluster - returns all results and visualization

    # Apply DR
    pca = PCA(n_components=n_components, random_state=42)
    X_pca = pca.fit_transform(X)

    ica = FastICA(n_components=n_components, random_state=42, max_iter=1000)
    X_ica = ica.fit_transform(X)"""

# =====#
# CELL 8: CLUSTERING AFTER DR (ALL 12 COMBINATIONS)
# =====#
print("\n" + "="*70)
print("STEP 3: CLUSTERING AFTER DIMENSIONALITY REDUCTION")
print("="*70)
```

```

rp = GaussianRandomProjection(n_components=n_components, random_state=42)
X_rp = rp.fit_transform(X)

dr_data = {'PCA': X_pca, 'ICA': X_ica, 'RP': X_rp}
results = {}

print(f"\n{dataset_name} Dataset (k={k}, n_components={n_components}):")
print("-" * 60)

for dr_name, X_t in dr_data.items():
    results[dr_name] = {}

    # K-Means
    km = KMeans(n_clusters=k, random_state=42, n_init=10)
    km_labels = km.fit_predict(X_t)
    km_sil = silhouette_score(X_t, km_labels)
    results[dr_name]['kmeans'] = {'labels': km_labels, 'silhouette': km_sil}

    # GMM
    gm = GaussianMixture(n_components=k, random_state=42, n_init=3)
    gm_labels = gm.fit_predict(X_t)
    gm_sil = silhouette_score(X_t, gm_labels)
    results[dr_name]['gmm'] = {'labels': gm_labels, 'silhouette': gm_sil}

    print(f" {dr_name}: K-Means={km_sil:.4f}, GMM={gm_sil:.4f}")

# Add to global results
all_results.append({'Dataset': dataset_name, 'DR_Method': dr_name,
                    'Clustering': 'K-Means', 'Silhouette': km_sil})
all_results.append({'Dataset': dataset_name, 'DR_Method': dr_name,
                    'Clustering': 'GMM', 'Silhouette': gm_sil})

# Visualization
fig, axes = plt.subplots(2, 3, figsize=(16, 10))

for idx, (dr_name, X_t) in enumerate(dr_data.items()):
    # K-Means
    sc = axes[0, idx].scatter(X_t[:, 0], X_t[:, 1],
                               c=results[dr_name]['kmeans']['labels'],
                               cmap='viridis', alpha=0.6, s=8)
    axes[0, idx].set_title(f'{dr_name} K-Means\nSilhouette: {results[dr_name]["kmeans"]["silhouette"]:.4f}')
    axes[0, idx].set_xlabel('Component 1')
    axes[0, idx].set_ylabel('Component 2')
    plt.colorbar(sc, ax=axes[0, idx])

    # GMM
    sc = axes[1, idx].scatter(X_t[:, 0], X_t[:, 1],
                               c=results[dr_name]['gmm']['labels'],
                               cmap='viridis', alpha=0.6, s=8)
    axes[1, idx].set_title(f'{dr_name} GMM\nSilhouette: {results[dr_name]["gmm"]["silhouette"]:.4f}')
    axes[1, idx].set_xlabel('Component 1')
    axes[1, idx].set_ylabel('Component 2')
    plt.colorbar(sc, ax=axes[1, idx])

```

```

plt.suptitle(f'{dataset_name}: Clustering After Dimensionality Reduction',
             fontsize=16, fontweight='bold')
plt.tight_layout()
plt.savefig(f'{dataset_name.lower()}_dr_clustering.png', dpi=150, bbox_inches='tight')
plt.show()

return results, dr_data

# Apply to BOTH datasets
cancer_dr_results, cancer_dr_data = dr_and_cluster(cancer_X_scaled, K, "Cancer")
bankrupt_dr_results, bankrupt_dr_data = dr_and_cluster(bankrupt_X_scaled, K,

```

---

### ===== STEP 3: CLUSTERING AFTER DIMENSIONALITY REDUCTION =====


Cancer Dataset (k=3, n\_components=2):

---

PCA: K-Means=0.3292, GMM=0.3294

ICA: K-Means=0.3378, GMM=0.3311

RP: K-Means=0.3600, GMM=0.3604




Bankruptcy Dataset (k=3, n\_components=2):

---

PCA: K-Means=0.3749, GMM=0.4081

ICA: K-Means=0.3428, GMM=0.4168

RP: K-Means=0.3152, GMM=0.4396



```
In [9]: # -----
# CELL 9: COMPREHENSIVE RESULTS TABLE
# -----
print("\n" + "="*70)
print("COMPREHENSIVE RESULTS TABLE")
print("="*70)

results_df = pd.DataFrame(all_results)
print(results_df.to_string(index=False))

# Summary statistics
print("\n--- Summary by Dataset and DR Method ---")
summary = results_df.groupby(['Dataset', 'DR_Method'])['Silhouette'].agg(['mean', 'count'])
print(summary)

# Best combinations
print("\n--- Top 5 Combinations by Silhouette Score ---")
print(results_df.nlargest(5, 'Silhouette').to_string(index=False))

# Save results
results_df.to_csv('clustering_results.csv', index=False)
print("\n✓ Results saved to clustering_results.csv")
```

=====

COMPREHENSIVE RESULTS TABLE

=====

Dataset	DR_Method	Clustering	Silhouette
Cancer	None	K-Means	0.086816
Cancer	None	GMM	0.072227
Bankruptcy	None	K-Means	0.107412
Bankruptcy	None	GMM	0.033795
Cancer	PCA	K-Means	0.329181
Cancer	PCA	GMM	0.329410
Cancer	ICA	K-Means	0.337764
Cancer	ICA	GMM	0.331061
Cancer	RP	K-Means	0.359981
Cancer	RP	GMM	0.360354
Bankruptcy	PCA	K-Means	0.374890
Bankruptcy	PCA	GMM	0.408054
Bankruptcy	ICA	K-Means	0.342803
Bankruptcy	ICA	GMM	0.416847
Bankruptcy	RP	K-Means	0.315174
Bankruptcy	RP	GMM	0.439565

--- Summary by Dataset and DR Method ---

mean      max

Dataset	DR_Method	mean	max
Bankruptcy	ICA	0.3798	0.4168
	None	0.0706	0.1074
	PCA	0.3915	0.4081
	RP	0.3774	0.4396
Cancer	ICA	0.3344	0.3378
	None	0.0795	0.0868
	PCA	0.3293	0.3294
	RP	0.3602	0.3604

--- Top 5 Combinations by Silhouette Score ---

Dataset	DR_Method	Clustering	Silhouette
Bankruptcy	RP	GMM	0.439565
Bankruptcy	ICA	GMM	0.416847
Bankruptcy	PCA	GMM	0.408054
Bankruptcy	PCA	K-Means	0.374890
Cancer	RP	GMM	0.360354

✓ Results saved to clustering\_results.csv

In [10]:

```
# =====
# CELL 10: STEP 4 - NEURAL NETWORK WITH DR
# =====
print("\n" + "="*70)
print("STEP 4: NEURAL NETWORK WITH DIMENSIONALITY REDUCTION")
print("="*70)

def nn_analysis(X, y, dataset_name, n_components=2):
    """NN performance with timing analysis"""

    le = LabelEncoder()
    y_enc = le.fit_transform(y)
```

```

X_train, X_test, y_train, y_test = train_test_split(
    X, y_enc, test_size=0.2, random_state=42, stratify=y_enc
)

nn_config = {
    'hidden_layer_sizes': (100, 50),
    'max_iter': 500,
    'random_state': 42,
    'early_stopping': True,
    'validation_fraction': 0.1
}

results = {}

print(f"\n{dataset_name} Neural Network Performance:")
print("-" * 60)

# Original
t0 = time.time()
mlp = MLPClassifier(**nn_config)
mlp.fit(X_train, y_train)
orig_time = time.time() - t0
orig_acc = mlp.score(X_test, y_test)
orig_f1 = f1_score(y_test, mlp.predict(X_test), average='weighted')

results['Original'] = {'accuracy': orig_acc, 'f1': orig_f1, 'time': orig_time}
print(f"  Original ({X.shape[1]} features):")
print(f"    Accuracy: {orig_acc:.4f}, F1: {orig_f1:.4f}, Time: {orig_time:.2f}s")

# DR methods
for dr_name, dr_model in [
    ('PCA', PCA(n_components=n_components, random_state=42)),
    ('ICA', FastICA(n_components=n_components, random_state=42, max_iter=1000)),
    ('RP', GaussianRandomProjection(n_components=n_components, random_state=42))
]:
    X_train_dr = dr_model.fit_transform(X_train)
    X_test_dr = dr_model.transform(X_test)

    t0 = time.time()
    mlp_dr = MLPClassifier(**nn_config)
    mlp_dr.fit(X_train_dr, y_train)
    dr_time = time.time() - t0
    dr_acc = mlp_dr.score(X_test_dr, y_test)
    dr_f1 = f1_score(y_test, mlp_dr.predict(X_test_dr), average='weighted')

    results[dr_name] = {'accuracy': dr_acc, 'f1': dr_f1, 'time': dr_time}
    print(f"\n  {dr_name} ({n_components} features):")
    print(f"    Accuracy: {dr_acc:.4f}, F1: {dr_f1:.4f}, Time: {dr_time:.2f}s")
    print(f"    Speedup: {(orig_time/dr_time):.2f}x")

# Visualization
fig, axes = plt.subplots(1, 3, figsize=(15, 5))

methods = list(results.keys())
accs = [results[m]['accuracy'] for m in methods]
f1s = [results[m]['f1'] for m in methods]

```

```

times = [results[m]['time'] for m in methods]
colors = ['#3498db', '#2ecc71', '#e74c3c', '#f39c12']

# Accuracy
bars = axes[0].bar(methods, accs, color=colors, alpha=0.8, edgecolor='black')
axes[0].set_ylabel('Test Accuracy', fontsize=12)
axes[0].set_title(f'{dataset_name}: Accuracy', fontsize=14, fontweight='bold')
axes[0].grid(True, alpha=0.3, axis='y')
for bar, acc in zip(bars, accs):
    axes[0].text(bar.get_x() + bar.get_width()/2, bar.get_height(),
                 f'{acc:.4f}', ha='center', va='bottom', fontsize=10, fontweight='bold')

# F1 Score
bars = axes[1].bar(methods, f1s, color=colors, alpha=0.8, edgecolor='black')
axes[1].set_ylabel('F1 Score', fontsize=12)
axes[1].set_title(f'{dataset_name}: F1 Score', fontsize=14, fontweight='bold')
axes[1].grid(True, alpha=0.3, axis='y')
for bar, f1 in zip(bars, f1s):
    axes[1].text(bar.get_x() + bar.get_width()/2, bar.get_height(),
                 f'{f1:.4f}', ha='center', va='bottom', fontsize=10, fontweight='bold')

# Time
bars = axes[2].bar(methods, times, color=colors, alpha=0.8, edgecolor='black')
axes[2].set_ylabel('Training Time (s)', fontsize=12)
axes[2].set_title(f'{dataset_name}: Training Time', fontsize=14, fontweight='bold')
axes[2].grid(True, alpha=0.3, axis='y')
for bar, t in zip(bars, times):
    axes[2].text(bar.get_x() + bar.get_width()/2, bar.get_height(),
                 f'{t:.2f}s', ha='center', va='bottom', fontsize=10, fontweight='bold')

plt.suptitle(f'{dataset_name}: Neural Network Performance', fontsize=16,
             fontweight='bold')
plt.tight_layout()
plt.savefig(f'{dataset_name.lower()}_nn_performance.png', dpi=150, bbox_inches='tight')
plt.show()

return results

# Run NN analysis
cancer_nn = nn_analysis(cancer_X_scaled, cancer_labels, "Cancer", n_components=2)
bankrupt_nn = nn_analysis(bankrupt_X_scaled, bankrupt_labels, "Bankruptcy", n_components=2)

```

---

---

#### STEP 4: NEURAL NETWORK WITH DIMENSIONALITY REDUCTION

---

---

##### Cancer Neural Network Performance:

---

Original (10 features):

Accuracy: 0.1273, F1: 0.1200, Time: 3.78s

PCA (2 features):


Accuracy: 0.1248, F1: 0.1055, Time: 2.86s  
Speedup: 1.32x

ICA (2 features):

Accuracy: 0.1270, F1: 0.1097, Time: 2.27s  
Speedup: 1.67x

RP (2 features):

Accuracy: 0.1219, F1: 0.1113, Time: 3.08s  
Speedup: 1.23x



##### Bankruptcy Neural Network Performance:

---

Original (96 features):

Accuracy: 0.9905, F1: 0.9897, Time: 0.43s

PCA (2 features):


Accuracy: 0.9685, F1: 0.9573, Time: 0.33s  
Speedup: 1.31x

ICA (2 features):

Accuracy: 0.9692, F1: 0.9588, Time: 0.55s  
Speedup: 0.78x

RP (2 features):

Accuracy: 0.9677, F1: 0.9519, Time: 0.26s  
Speedup: 1.66x



In [11]:

```
# =====
# CELL 11: HYPOTHESIS VALIDATION
# =====

print("\n" + "="*70)
print("HYPOTHESIS VALIDATION")
print("="*70)

# Hypothesis 1
print("\nHYPOTHESIS 1: DR improves clustering quality")
print("-" * 60)

for ds in ['Cancer', 'Bankruptcy']:
    orig = results_df[(results_df['Dataset'] == ds) & (results_df['DR_Method'] == 'Original')]
    after_dr = results_df[(results_df['Dataset'] == ds) & (results_df['DR_Method'] == 'DR')]
    improvement = ((after_dr - orig) / orig) * 100

    print(f"\n{ds}:")
    print(f" Original: {orig:.4f}")
    print(f" After DR: {after_dr:.4f}")
    print(f" Improvement: {improvement:+.1f}%")
    print(f" Result: {'✓ CONFIRMED' if after_dr > orig else '✗ REJECTED'}")

# Hypothesis 2
print("\nHYPOTHESIS 2: RP performs comparably to PCA/ICA")
print("-" * 60)

for ds in ['Cancer', 'Bankruptcy']:
    pca = results_df[(results_df['Dataset'] == ds) & (results_df['DR_Method'] == 'PCA')]
    ica = results_df[(results_df['Dataset'] == ds) & (results_df['DR_Method'] == 'ICA')]
    rp = results_df[(results_df['Dataset'] == ds) & (results_df['DR_Method'] == 'RP')]

    diff = abs(rp - (pca + ica)/2) / ((pca + ica)/2) * 100

    print(f"\n{ds}:")
    print(f" PCA: {pca:.4f}, ICA: {ica:.4f}, RP: {rp:.4f}")
    print(f" RP differs by: {diff:.1f}%")
    print(f" Result: {'✓ CONFIRMED' if diff < 15 else '~ PARTIAL'}")

# Hypothesis 3
print("\nHYPOTHESIS 3: Original space clustering less meaningful")
print("-" * 60)

for ds in ['Cancer', 'Bankruptcy']:
```

```
orig = results_df[(results_df['Dataset'] == ds) & (results_df['DR_Method'] == 'original')]
best = results_df[(results_df['Dataset'] == ds) & (results_df['DR_Method'] == 'best')]

print(f"\n{ds}:")
print(f"  Original: {orig:.4f}")
print(f"  Best DR: {best:.4f}")
print(f"  Result: {'✓ CONFIRMED' if orig < 0.2 else '~ PARTIAL'}")

print("\n" + "="*70)
print("ANALYSIS COMPLETE")
print("="*70)
print("\nGenerated files:")
print("  - cancer_optimal_k.png")
print("  - bankruptcy_optimal_k.png")
print("  - cancer_pca_scree.png")
print("  - bankruptcy_pca_scree.png")
print("  - cancer_ica_kurtosis.png")
print("  - bankruptcy_ica_kurtosis.png")
print("  - cancer_rp_reconstruction.png")
print("  - bankruptcy_rp_reconstruction.png")
print("  - cancer_dr_clustering.png")
print("  - bankruptcy_dr_clustering.png")
print("  - cancer_nn_performance.png")
print("  - bankruptcy_nn_performance.png")
print("  - clustering_results.csv")
```

---

---

## HYPOTHESIS VALIDATION

---

---

---

### HYPOTHESIS 1: DR improves clustering quality

---

Cancer:

Original: 0.0795  
After DR: 0.3413  
Improvement: +329.2%  
Result: ✓ CONFIRMED

Bankruptcy:

Original: 0.0706  
After DR: 0.3829  
Improvement: +442.3%  
Result: ✓ CONFIRMED

---

### HYPOTHESIS 2: RP performs comparably to PCA/ICA

---

Cancer:

PCA: 0.3293, ICA: 0.3344, RP: 0.3602  
RP differs by: 8.5%  
Result: ✓ CONFIRMED

Bankruptcy:

PCA: 0.3915, ICA: 0.3798, RP: 0.3774  
RP differs by: 2.1%  
Result: ✓ CONFIRMED

---

### HYPOTHESIS 3: Original space clustering less meaningful

---

Cancer:

Original: 0.0795  
Best DR: 0.3604  
Result: ✓ CONFIRMED

Bankruptcy:

Original: 0.0706  
Best DR: 0.4396  
Result: ✓ CONFIRMED

---

---

ANALYSIS COMPLETE

---

---

Generated files:

- cancer\_optimal\_k.png
- bankruptcy\_optimal\_k.png
- cancer\_pca\_scree.png
- bankruptcy\_pca\_scree.png

```
- cancer_ica_kurtosis.png  
- bankruptcy_ica_kurtosis.png  
- cancer_rp_reconstruction.png  
- bankruptcy_rp_reconstruction.png  
- cancer_dr_clustering.png  
- bankruptcy_dr_clustering.png  
- cancer_nn_performance.png  
- bankruptcy_nn_performance.png  
- clustering_results.csv
```

In [ ]: